

Laboratory Report

Core-affine threaded interpolating elevation of a $n \times n$ matrix M given a lower resolution digital elevation matrix N

Jasrel Roby Peralta

May 2023

Introduction

After creating a threaded computer program from Exercise 02, use the programming exercise from Exercise 03 to record average runtimes of estimation of a $n \times n$ matrix when using a different t number of threads.

Objectives

The goal for this exercise is the following:

- determine the complexity of estimating the point elevation of a $n \times n$ square matrix with randomized values at grid points divisible by 10 when using n concurrent processors and other values of concurrent processors.
- see if the runtime for this exercise is lower than the average runtime that was obtained in Exercise 01 and Exercise 02.
- figure out why higher values of n size of matrix are now faster using t concurrent cores.

Methodology

The machine used for this exercise is running on Ubuntu 22.04.2 LTS x86_64, Intel i-7 8700 (12 cores) @ 4.60GHz, AMD ATI Radeon HD 8570 / RS 430, with 16GB memory. The programming language used in the computer program is Python 3.10.6. The interpolating algorithm used was the Federal Communications Commission (FCC) method. The graphing software used for making the charts is LibreOffice Calc.

The computer program made use of the *multiprocessing* module of Python 3.10.6, to utilize t cores and concurrently estimate different $(n/t) \times n$ submatrices from the $n \times n$ matrix.

The size of the matrix for all recorded runs was $n = 8000$. Moreover, three (3) runs were done using t number of threads, starting from 1 (2^0) up to 64 (2^6). The created threads are then automatically assigned by Python to different cores. These runs were then averaged and recorded to a table. The program was further analyzed using *htop*, an interactive process viewer, to see if the cores being assigned are being used to its full capacity.

Results and Discussion

After running the code three (3) successive times, using n matrix size and t cores, the following table is produced:

n (size of matrix)	t (number of concurrent cores)	Time Elapsed (seconds)			Average Runtime (seconds)
		Run 1	Run 2	Run 3	
8000	1	99.163350	97.724076	97.501831	98.129752
8000	2	51.743841	52.954694	49.251313	51.316616
8000	4	27.249401	26.473527	26.383961	26.702296
8000	8	20.739265	21.435630	22.115806	21.430234
8000	16	20.819450	20.769891	20.854711	20.814684
8000	32	21.140523	20.939167	21.187100	21.088930
8000	64	22.781048	20.981064	21.185897	21.649336

Table 1: Average runtimes of the computer program with 1 to 64 threads

To further understand the table, the following line chart is created using LibreOffice:




Figure 1: Line Chart of the Average runtimes of the computer program with 1 to 64 threads (multiprocessing)

As seen in Table 1, the running times of the code using different number of t cores. As the number of cores increase, the runtimes decrease. This progression can be easily observed during the increasing of number of t cores while dealing with low values, but as the number of cores increase, the decrease in runtime flattens as it continues.

To compare the recently obtained data to the previously obtained data using multithreading in the previous exercise, the table from the previous laboratory report is shown below:

n (size of matrix)	t (number of concurrent threads)	Time Elapsed (seconds)			Average Runtime (secs)
		Run 1	Run 2	Run 3	
8000	1	100.912083	99.045359	98.716483	99.557975
8000	2	106.021552	107.055136	105.227246	106.101311333333
8000	4	108.299674	107.104734	107.16318	107.522529333333
8000	8	114.795173	113.608294	122.527366	116.976944333333
8000	16	119.669082	118.507086	119.356774	119.177647333333
8000	32	124.609196	124.346253	122.94652	123.967323
8000	64	125.016126	130.814412	125.503346	127.111294666667

Table 2: Average runtimes of the computer program from 1 to 64 concurrent threads (multithreading)

In comparing Tables 1 and 2, it is observed that the trend in both tables are different. The table obtained in Table 2 shows an inverse relationship with the n size of matrix and t number of threads, while Table 1 displays a direct relationship with the n size of matrix and t number of threads.

The reason for the difference in trends in tables of multiprocessing (Table 1) and multithreading (Table 2) is the Global Interpreter Lock (GIL) present when using the *multithreading* module [1, 2]. GIL is being applied when dealing with mulple threads in Python, but not when using multiple processes. Hence, the decrease in runtime in multiprocessing, unlike in multithreading [3].




Figure 2: Screen capture of *htop* during the execution of the computer program using a single thread

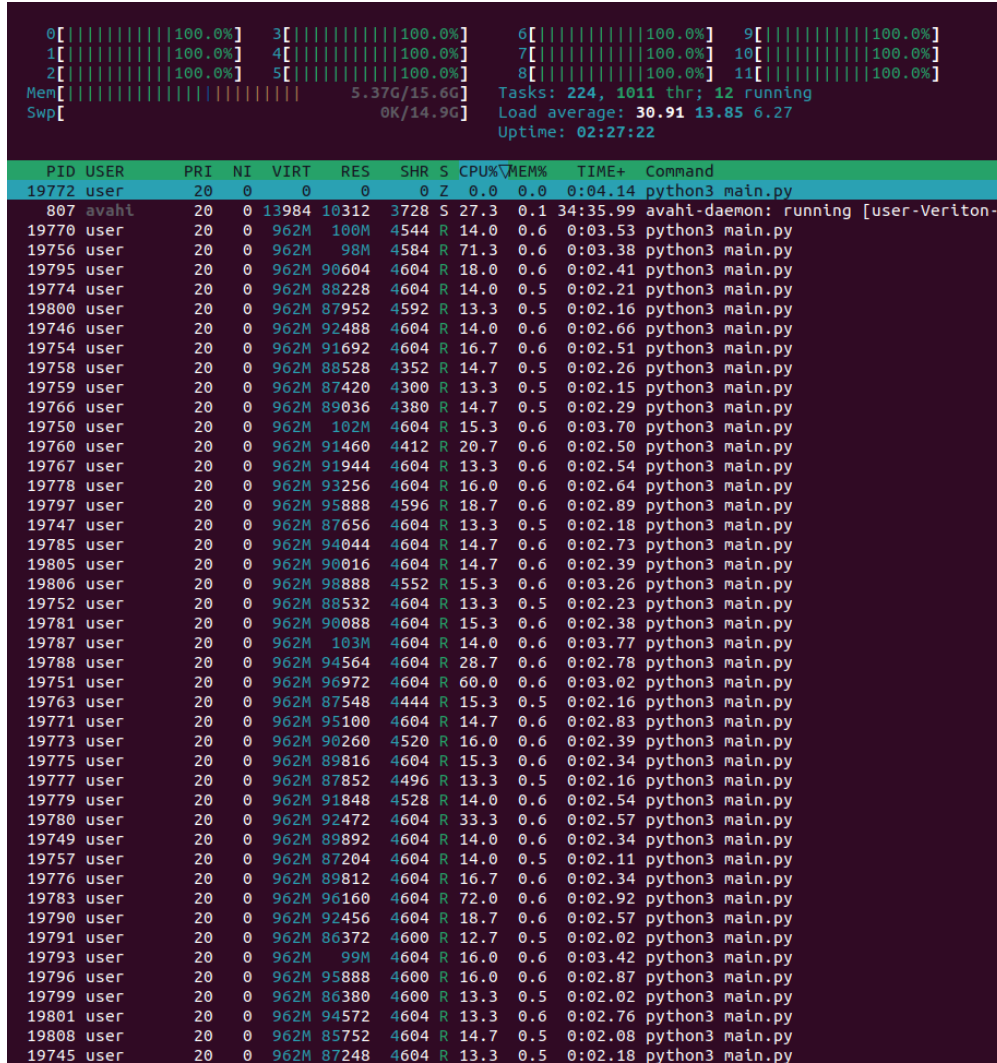


Figure 3: Screen capture of *htop* during the execution of the computer program using 64 threads

Figures 2 and 3 shows the screen captures of *htop* while running the program. The CPU% column shows what percentage of the capacity of the memory is being utilized by the process [4]. In Figure 2, it is seen that a single core is being maximized. And in Figure 3, all cores are being maximized during runtime.

Conclusion

The complexity of estimating the point elevation of a $n \times n$ square matrix with randomized values at grid points divisible by 10 when using n concurrent threads is $O(n)$ because each column is iterated through only once, since each column will be assigned to a single thread, which will be then be assigned to a core. This complexity is still the same with the multithreading program, but since GIL is not being utilized in the *multiprocessing* module of Python, the speedup is observed.

Compared to Exercise 01 and 02, the runtimes obtained in this exercise is much faster. This also made it possible for higher $n \times n$ square matrices to be possible to be interpolated since the GIL is now hindering the processes in running concurrently.

References

- [1] Real Python. *An Intro to threading in Python*. 2022. URL: <https://realpython.com/intro-to-python-threading>. [Accessed: 28-Mar-2023].
- [2] Real Python. *What is the python global interpreter lock (gil)?* 2021. URL: <https://realpython.com/python-gil>. [Accessed: 28-Mar-2023].
- [3] Marcus McCurdy. *Python Multithreading and Multiprocessing Tutorial*. 2023. URL: <https://www.toptal.com/python/beginners-guide-to-concurrency-and-parallelism-in-python>. [Accessed: 2-May-2023].
- [4] Michael Kerrisk. *htop(1) — Linux manual page*. 2022. URL: <https://man7.org/linux/man-pages/man1/htop.1.html>. [Accessed: 2-May-2023].

Appendix

Interpolation Source Code

(main.py)

```
import numpy as np
import random
import datetime
import os
import threading
from multiprocessing import Process

print("This machine has", os.cpu_count(), "number of CPUs")

# prettier printing options
np.set_printoptions(linewidth=1000, formatter={'float': '{:0.0f}'.format})

# interpolate function
def terrain_inter_multithreading(mat, x1, x2):
    for i in range(0, n):
        for j in range(x1, x2):
            if mat[i][j] != 0:
                continue
            if (i % dist == 0):
                get_row_val(i, j)
    for i in range(0, n):
        for j in range(x1, x2):
            if (mat[i][j] == 0):
                get_col_val(i, j)

# modified interpolation function
# to run concurrently with other threads
# from other x1 to x2's
def terrain_inter_multiprocessing(mat, x1, x2):
    for i in range(0, n):
        for j in range(x1, x2+2):
            if mat[i][j] != 0:
                continue
            if (i % dist == 0):
                get_row_val(i, j)
    for i in range(0, n):
        for j in range(x1, x2+2):
            if (mat[i][j] == 0):
                get_col_val(i, j)

def get_submatrices(n, t):
    # array of submatrices
    sub_arr = []
    temp = []
    for i in range(0, n):
```

```

        temp.append(i)
        if (len(temp) == (n-1) / t):
            sub_arr.append(temp)
            temp = []

    return sub_arr

# get size of matrix
def getSize():
    n = 1
    while (n % 10 != 0):
        n = int(input("enter_size_of_matrix:_"))
        if n % 10 != 0:
            print('invalid_size_of_matrix')
    return n+1

# get number of threads
def getThreads(n):
    n -= 1
    t = 0
    # n size should be less than t threads
    # t threads should not be 0
    # n size should be divisible by t threads
    while (n < t) or (t == 0) or (n % t != 0):
        t = int(input('enter_number_of_threads:_'))
        if (n < t) or (n % t != 0):
            print('invalid_number_of_threads')
    return t

def getCores(n):
    n -= 1
    t = 0
    # n size should be less than t threads
    # t threads should not be 0
    # n size should be divisible by t threads
    while (n < t) or (t == 0) or (n % t != 0):
        t = int(input('enter_number_of_threads:_'))
        if (n < t) or (n % t != 0):
            print('invalid_number_of_threads')
    return t

# dp array format:
# dp = [[x1,y1][x2,y2]]

# interpolate rows with random values
def get_row_val(i,j):
    dp = get_datapoints_row(i,j)
    x = j # j -> row
    x1 = dp[0][0]
    x2 = dp[1][0]
    y1 = dp[0][1]
    y2 = dp[1][1]
    res = fcc(x1,y1,x2,y2,x)
    mat[i][j] = res

```

```

# interpolate columns
def get_col_val(i,j):
    dp = get_datapoints_col(i,j)
    # dp = [[x1,y1][x2,y2]]
    x = i # i -> col
    x1 = dp[0][0]
    x2 = dp[1][0]
    y1 = dp[0][1]
    y2 = dp[1][1]
    res = fcc(x1,y1,x2,y2,x)
    mat[i][j] = res

# get closest datapoints to the current gridpoint
def get_datapoints_row(i,j):
    dp = []
    dp.append(get_nearest_row(i,j,-1))
    dp.append(get_nearest_row(i,j,+1))
    return dp
def get_datapoints_col(i,j):
    dp = []
    dp.append(get_nearest_col(i,j,-1))
    dp.append(get_nearest_col(i,j,+1))
    return dp

# x, y -> point; dir -> direction
# change direction to check to the nearest 10
## improved from recursion from previous exercise to direct computation
def get_nearest_row(i,j,dir):
    # go up
    if dir < 0:
        dir = j - (j % 10)
    # go down
    else:
        dir = j + (10 - (j % 10))
    return [dir,mat[i][dir]]

def get_nearest_col(i,j,dir):
    # go left
    if dir < 0:
        dir = i - (i % 10)
    # go right
    else:
        dir = i + (10 - (i % 10))
    return [dir,mat[dir][j]]

# follow given FCC formula
def fcc(x1,y1,x2,y2,x):
    return (y1 + (((x-x1)/(x2-x1)) * (y2-y1)))

# main function
if __name__ == "__main__":
    # initialize data

```



```

n = getSize()
t = getCores(n)

# distance between randomized values
dist = 10

# create a zero nxn matrix
mat = np.zeros((n,n), dtype = float)

# randomize elevation values for gridpoints divisible by 10
for i in range(n):
    for j in range(n):
        if i % dist == 0 and j % dist == 0:
            mat[i][j] = random.uniform(0.0, 1000.0)

# print initial matrix
print(mat)

threads = list()

for set in get_submatrices(n,t):
    x1, x2 = set[0], set[-1]
    thread = threading.Thread(target=terrain_inter_multithreading, args=(mat,x1,x2))
    threads.append(thread)

# record time before threaded interpolation
time_before_multithreading = datetime.datetime.now()

for thread in threads:
    thread.start()

for thread in threads:
    thread.join()

# record time after threaded interpolation
time_after_multithreading = datetime.datetime.now()
# print resulting matrix
print(mat)

print("\n\n\n")

# create a zero nxn matrix
mat = np.zeros((n,n), dtype = float)

# randomize elevation values for gridpoints divisible by 10
for i in range(n):
    for j in range(n):
        if i % dist == 0 and j % dist == 0:
            mat[i][j] = random.uniform(0.0, 1000.0)

# print initial matrix
print(mat)

```

```

processes = list()

for set in get_submatrices(n,t):
    x1, x2 = set[0], set[-1]
    process = Process(target=terrain_inter_multiprocessing, args=(mat,x1,x2))
    processes.append(process)

# record time before threaded interpolation
time_before_multiprocessing = datetime.datetime.now()

for process in processes:
    process.start()

for process in processes:
    process.join()

# record time after threaded interpolation
time_after_multiprocessing = datetime.datetime.now()

# print resulting matrix
print(mat)

# print interpolation time
print("multithreading: ", time_after_serial - time_before_serial)
print("multiprocessing: ", time_after_multiprocessing - time_before_multiprocessing)

```